

---

# *Computer Architecture*

## Chapter 1: Performance

**Joonho Kong**

**School of Electronics Engineering  
Kyungpook National University**

\* Thanks Prof. Sung Woo Chung in Korea University allowing me to use the ppt materials

# Main points for this chapter

---

- How can we quantify performance?
- What are the metrics for performance?
- To improve performance, how can we do?
- What factors do affect performance?
- What another metric for performance and their pitfalls?
- What is benchmark?
- What are Amdahl's law and its implications for computer architecture design?
- What is power wall?

# How to define performance

---

- There are many ways to define something as “the best”
- Airplane example

Airplane	Passenger Capacity	Cruising time (miles)	Cruising speed (m.p.h)	Throughput (passengers * m.p.h)
Boeing 777	375	5,256	610	228,750
Boeing 747	① 416	7,156	610	286,700
Airbus 380	525	8,200	560	① 294,000
BAC/Sud Concorde	132	4,000	① 1,350	178,200
Douglas DC-8-50	146	① 8,720	544	79,424

- What is “the best”?

# Applying to computers...

---

- How do you decide which computer is the best?
  - Processor speed
  - System bus speed
  - Memory size
  - Disk storage
  - Graphics card / subsystem
  - Power consumption
  - Price
- How do you decide which one to buy?
  - Tradeoff between cost and performance

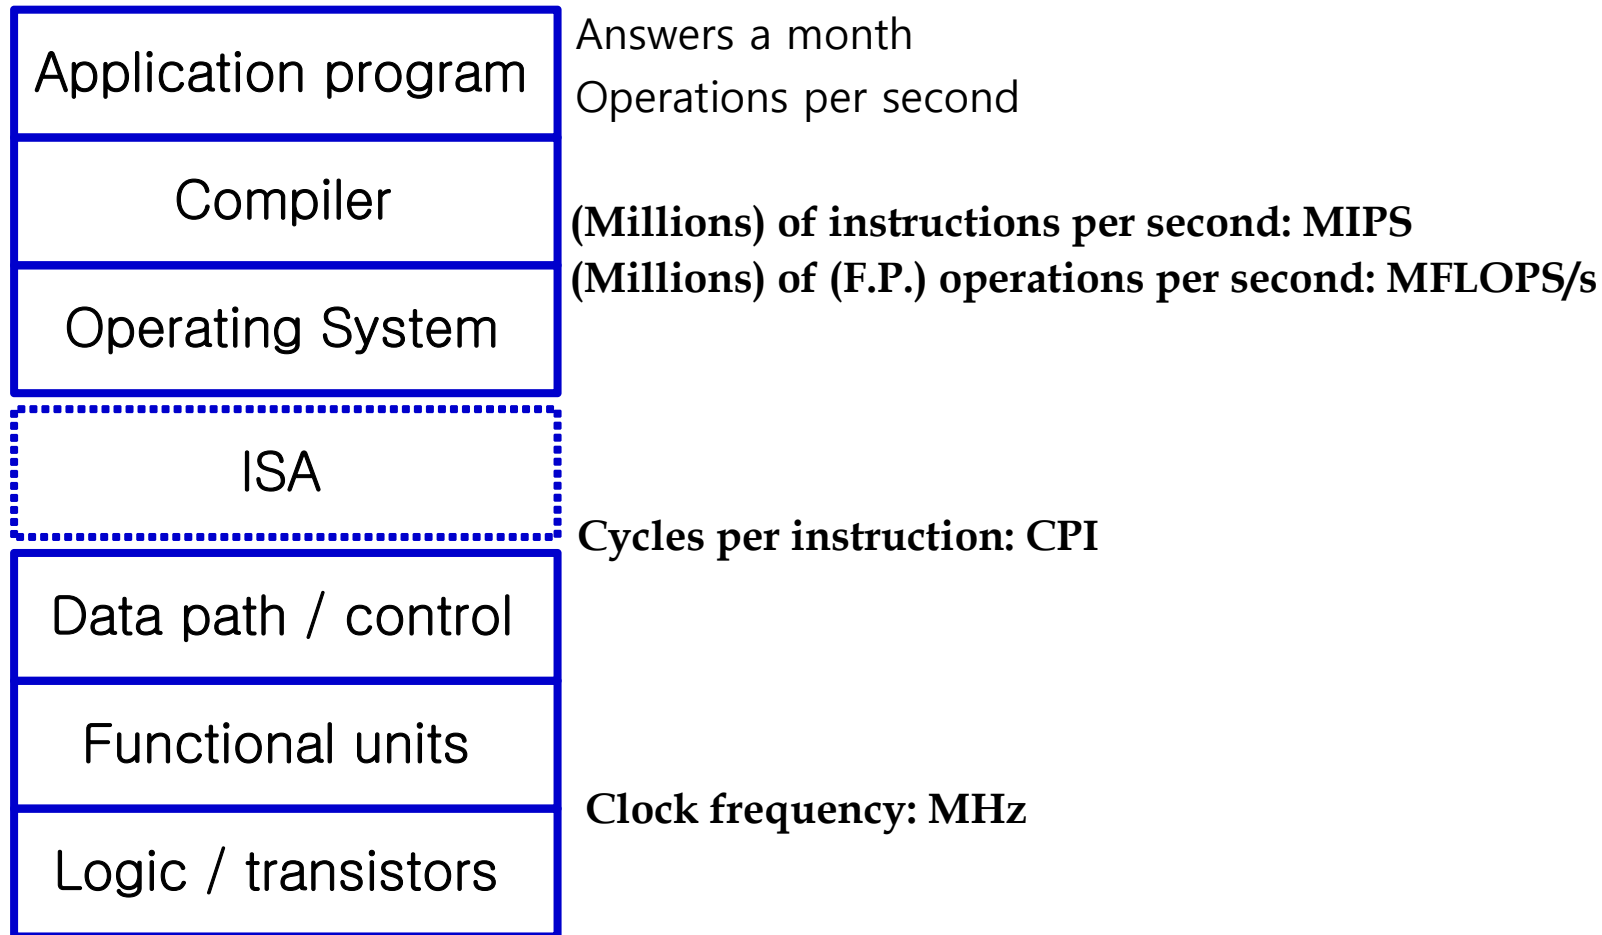
# Quantifying the performance

---

- Important when
  - Purchasing a computer
  - Evaluating new technologies
  - Writing software
  - Implementing an instruction set
  - Designing a new architecture
- Therefore, it is important to understand how to define performance and what the limitations are of those metrics

# Example of metrics in computers

---



- ❑ Most of metrics are related to time
  - Time is the most classical metric for computers

# Performance metrics for computer systems

---

- Execution time
  - How fast?
  - time taken for doing a certain work
  - e.g., This computer takes 10 seconds to run a certain program
- Throughput
  - How much?
  - amount of work done per time
  - e.g., This SSD can read or write 10 Giga-byte data per second

# Measuring time

---

- Looking at measuring CPU performance, we are primarily concerned with execution time

$$\text{Performance} = \frac{1}{\text{Execution time}}$$

- To compare, we say “X is n times faster than Y”

$$n = \frac{\text{Performance}_x}{\text{Performance}_y} = \frac{\text{Execution time}_y}{\text{Execution time}_x}$$

- Increase performance, decrease execution time
  - Improve performance, improve execution time



# Example of performance comparison

---

- X and Y do their homework
  - X takes 5 hours
  - Y takes 10 hours
- Compare the performance

$$\text{Performance}_x = \frac{1}{\text{Execution time}_x} = \frac{1}{5 \text{ hours}} = 0.2$$

$$\text{Performance}_y = \frac{1}{\text{Execution time}_y} = \frac{1}{10 \text{ hours}} = 0.1$$

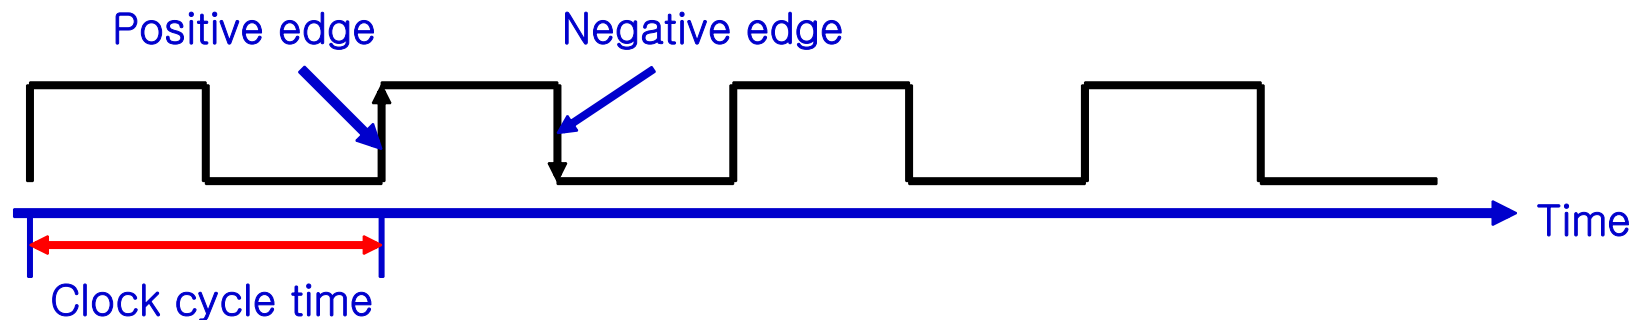
- So, X is two times faster than Y

$$n = \frac{\text{Performance}_x}{\text{Performance}_y} = \frac{0.2}{0.1} = 2$$

# Clock cycle time vs. Clock rate

---

- Clock cycle time
  - Time required for a clock pulse to make transitions:  $0 \rightarrow 1 \rightarrow 0$

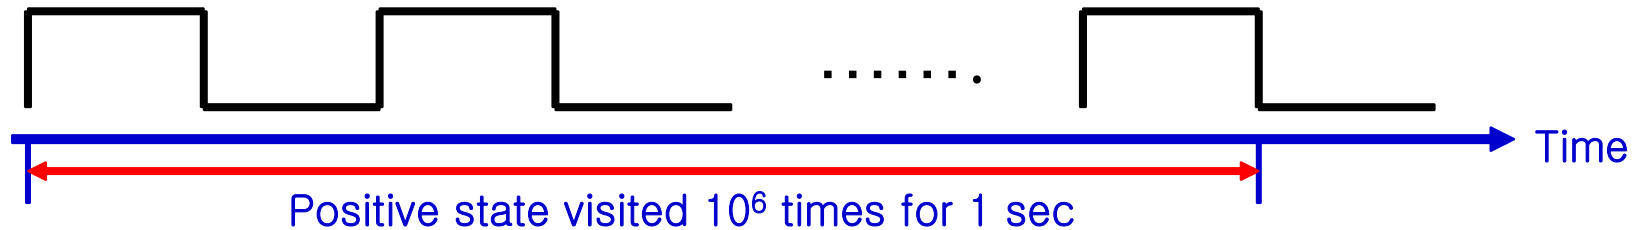


- In other words, the time duration between positive (negative) edges
- Clock rate
  - Inverse of clock cycle time
  - # of times to visit positive (negative) state per second
  - Unit: Hz or MHz or GHz

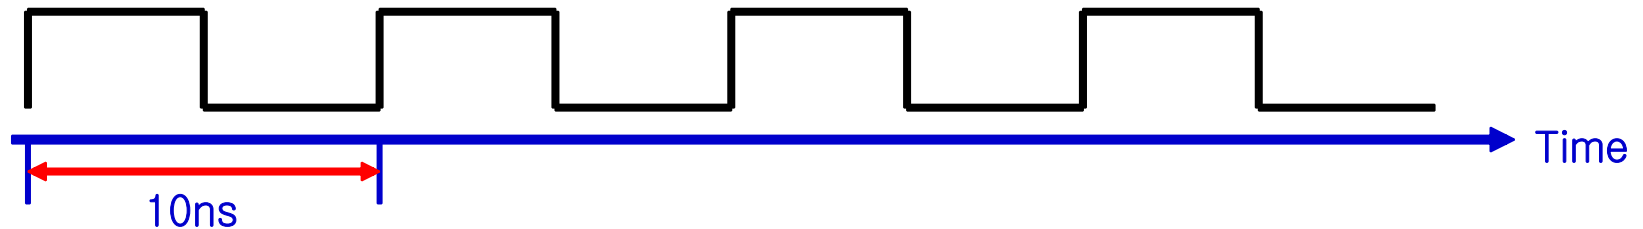
# Example of clock cycle time

- A machine is running at 100MHz

- Clock rate = 100MHz =  $100 * 10^6$  cycles / sec

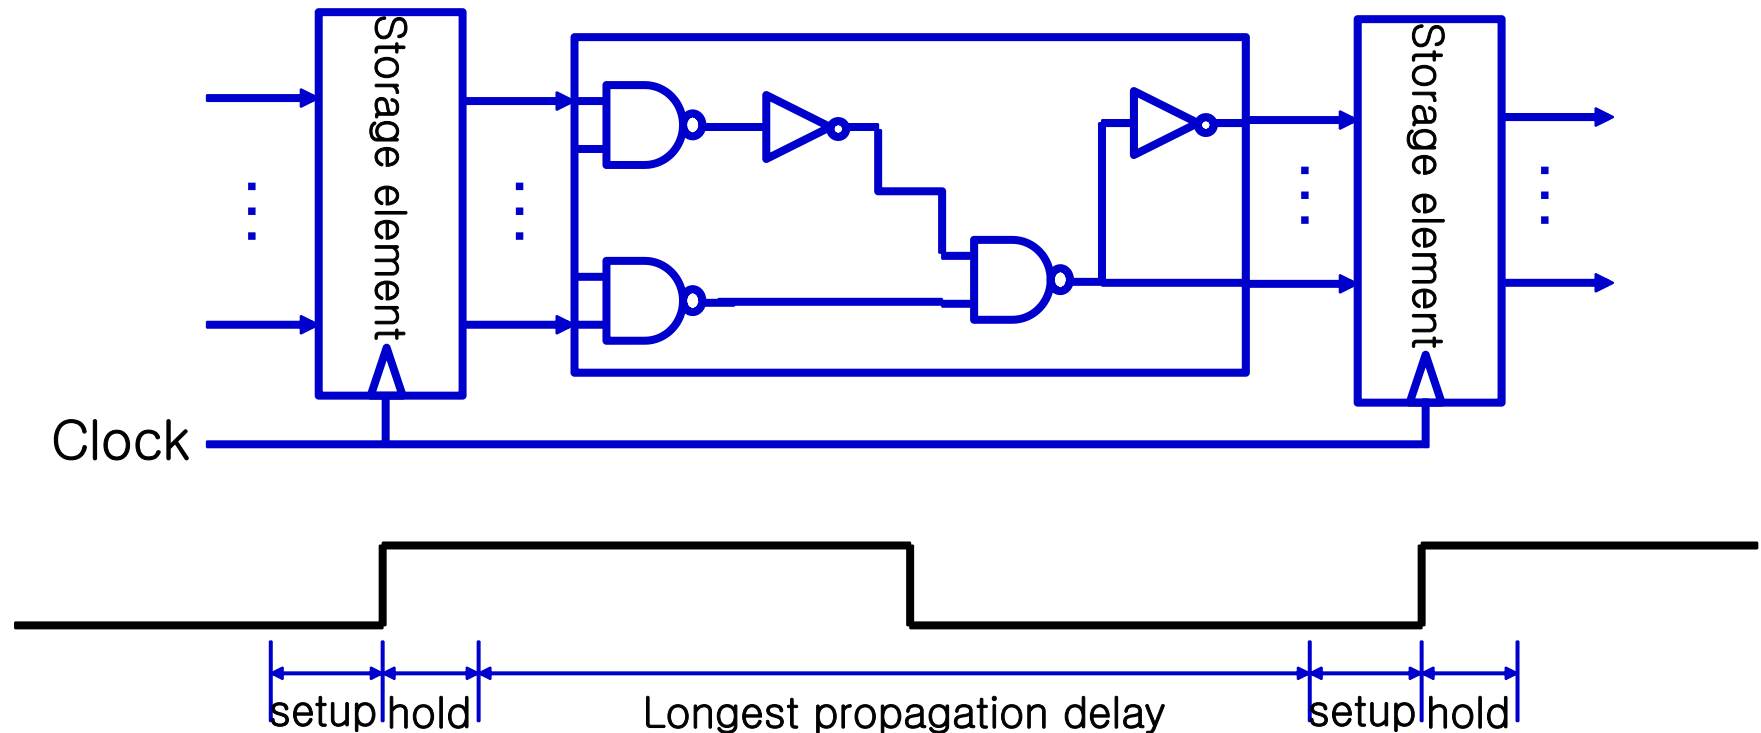


- Clock cycle time =  $1/(100*10^6)$  cycles / sec = 10ns



# How is clock cycle time determined?

- Closely related to logic design
- Assumptions
  - All storage elements have same clock
  - Edge-triggered clocking
  - Design always work if the clock is slow enough



# Meaning of timing parameters

---

- Longest propagation delay
  - A.k.a critical path delay
  - Critical path: a path that takes the most timing delays among many combinational paths
  - Typically identified by timing analysis or static timing analysis

# Execution time

---

- We will use CPU execution time frequently as the metric of how long a program should run
  - Execution time = Clock cycles for program \* Clock cycle time
- Since clock cycle time is the inverse of clock rate
  - **Execution time =  $\frac{\text{Clock cycles for program}}{\text{Clock rate}}$**

# Measuring clock cycles

---

- CPU clock cycles / program is not so intuitive
  - Program is a set (ordered) of instructions
- CPI (Cycles Per Instruction) is used so frequently
  - The # of cycles per instruction varies, so CPI is an average value
  - IPC?

# Using CPI

---

- Therefore, we can rewrite
  - Execution time = Instructions \* CPI \* Clock cycle time
  - Improved performance (reduced execution time) is possible with increased clock rate (reduced clock cycle time), lower CPI, or reduced instructions
  - Designers have to balance the length of each cycle and the number of cycles required



# CPI Example

---

- Machine A: 1ns clock and CPI of 2.0
- Machine B: 2ns clock and CPI of 0.5
- Which is faster?

# Example solution

---

- Solve CPU time for each machine
  - Execution time<sub>A</sub> = I \* 2.0 \* 1ns = 2.0 \* I ns
  - Execution time<sub>B</sub> = I \* 0.5 \* 2ns = 1.0 \* I ns
- Compare performance

$$\frac{\text{Performance}_A}{\text{Performance}_B} = \frac{\text{Execution time}_B}{\text{Execution time}_A} = \frac{1.0 * I \text{ ns}}{2.0 * I \text{ ns}} = 0.5$$

- So, machine A is 0.5 times faster than machine B  
= Machine B is 2 times faster than machine A
  - You must consider both
    - CPI
    - Clock rate

# CPI variability

---

- Different types of instructions often take different numbers of cycles on the same processor
- CPI is often reported for classes of instructions
  - Clock cycles =  $\sum_{i=1}^n ( \mathbf{CPI}_i \times \mathbf{C}_i )$
  - $\mathbf{CPI}_i$  : the CPI for the class of instructions
  - $\mathbf{C}_i$ : the count of that type of instructions

# CPI from instruction mix

- $$\text{CPI} = \frac{\sum_{i=1}^n (\text{CPI}_i \times C_i)}{\text{Instruction Count}} = \sum_{i=1}^n (\text{CPI}_i \times \frac{C_i}{\text{Instruction Count}})$$

$\frac{C_i}{\text{Instruction Count}}$

frequency of appearance of the type i instructions
- CPI Example

Instruction Class	Appearance Frequency	CPI <sub>i</sub>
Instruction type 1	43%	1
Instruction type 2	21%	2
Instruction type 3	12%	2
Instruction type 4	24%	2

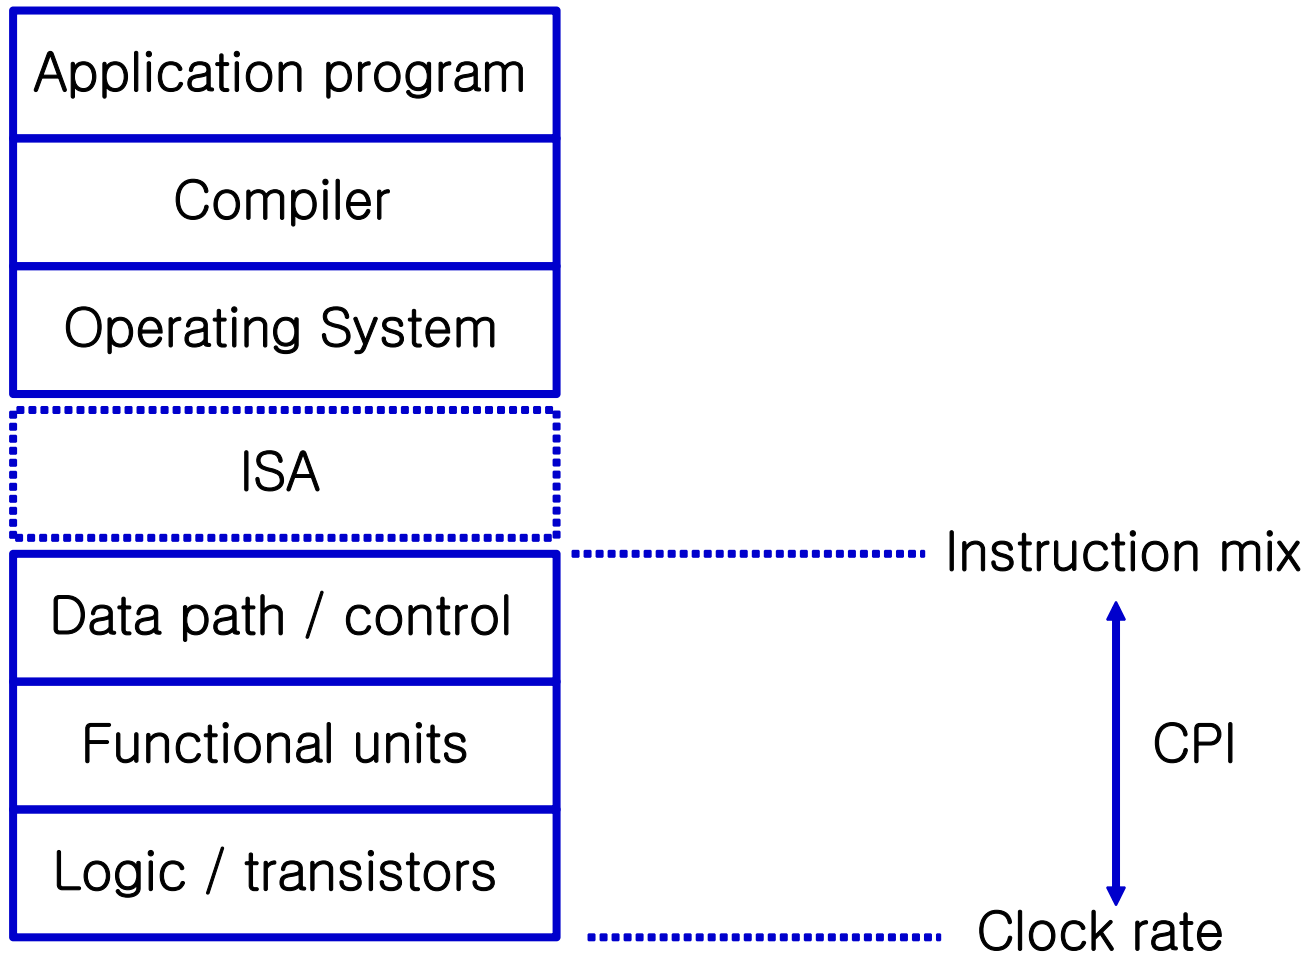
$$\text{CPI} = 0.43 \times 1 + 0.21 \times 2 + 0.12 \times 2 + 0.24 \times 2 = 1.57$$

$$\text{Clock cycles} = 1.57 * \text{Instruction Count}$$

# Tradeoffs

---

- Instruction count, CPI, and clock rate present tradeoffs



# Aspects of CPU performance

$$\text{CPU time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

	Instruction count	CPI	Clock rate
Program	✓		
Compiler	✓		
ISA	✓	✓	
Organization		✓	✓
Technology			✓

□ Ambiguous!! You only have to understand the meaning

# Another popular performance metrics

---

- MIPS (million instructions per second)
  - $\text{MIPS} = \frac{\text{Instruction count}}{\text{Execution time} \times 10^6}$
  - Problems
    - It does not take into account the instruction set
- MFLOPS (million floating-point operations per second)
  - Operations rather than instruction
  - E.g. floating-point addition, multiplication, ...
  - Often used for quantifying supercomputing performance or domain-specific architecture performance

# A wrong use case of MIPS

---

- Consider a 500MHz machine

Class	CPI
Class A	1
Class B	2
Class C	3

- Consider the two compilers

Code from	Instruction counts (millions)		
	A	B	C
Compiler1	5	1	1
Compiler2	10	1	1

- Which compiler produce faster code? Has a higher MIPS?



# A wrong use case of MIPS: solution (I)

---

- Compute clock cycles
  - Clock cycles =  $\sum_{i=1}^n (CPI_i \times C_i)$
  - Clock cycles<sub>comp1</sub> =  $(1 \times 5M) + (2 \times 1M) + (3 \times 1M) = 10M$
  - Clock cycles<sub>comp2</sub> =  $(1 \times 10M) + (2 \times 1M) + (3 \times 1M) = 15M$
- Execution time
  - Execution time = (Instruction count \* CPI) / Clock rate  
= Clock cycles / Clock rate
  - Execution time<sub>comp1</sub> =  $10M / 500M = 0.02\text{sec}$
  - Execution time<sub>comp2</sub> =  $15M / 500M = 0.03\text{sec}$
- Code from compiler 1 is 1.5 times faster!

# A wrong use case of MIPS: solution (II)

---

- Computer MIPS
  - $$\text{MIPS} = \frac{\text{Instruction count}}{\text{Execution time} \times 10^6}$$
  - $\text{MIPS}_{\text{comp1}} = (5\text{M} + 1\text{M} + 1\text{M}) / (0.02\text{M}) = 350$
  - $\text{MIPS}_{\text{comp2}} = (10\text{M} + 1\text{M} + 1\text{M}) / (0.03\text{M}) = 400$
- Code from compiler 2 is faster??
  - Fails to give a right answer!

# Benchmarks

---

- Users often want a performance metric
- A benchmark is distillation of the attributes of a workload
  - Real applications usually work best, but using them is not always feasible
- Desirable attributes
  - Relevant: meaningful within the target domain
  - Understandable
  - Good metric(s)
  - Scalable
  - Coverage: does not oversimplify important factors in the target domain
  - Acceptance: vendors and users embrace it

# Benchmarks (cont)

---

- de facto industry standard benchmarks for CPU
  - SPEC
- SPEC
  - Standard Performance Evaluation Cooperative
  - Founded in 1988 by EE times, SUN, HP, MIPS, Apollo, DEC
  - Several different SPEC benchmarks
  - Most include a suite of several different applications (such as integer and floating point components often reported separately)
  - For more information, visit <http://www.spec.org>
  - The latest version: SPEC CPU 2017
- Synthetic benchmarks do not come from real executables, but are designed to approximate a machine's performance
  - Synthetic benchmarks such as Whetstone and Dhrystone are not really measuring anything people run and can be optimized

# Amdahl's law

---

Execution speedup is proportional to the size of the improvement and the amount affected

- Execution time after improvement

$$= \left[ \frac{\text{Execution time affected by improvement}}{\text{Amount of improvement}} + \text{Execution time unaffected} \right]$$

- Or

$$\text{ExTime}_{\text{new}} = \text{ExTime}_{\text{old}} \times \left[ (1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}} \right]$$

# Example – Amdahl's law

---

- Floating point instructions improved to run 2X; but only 10% of actual instructions are FP

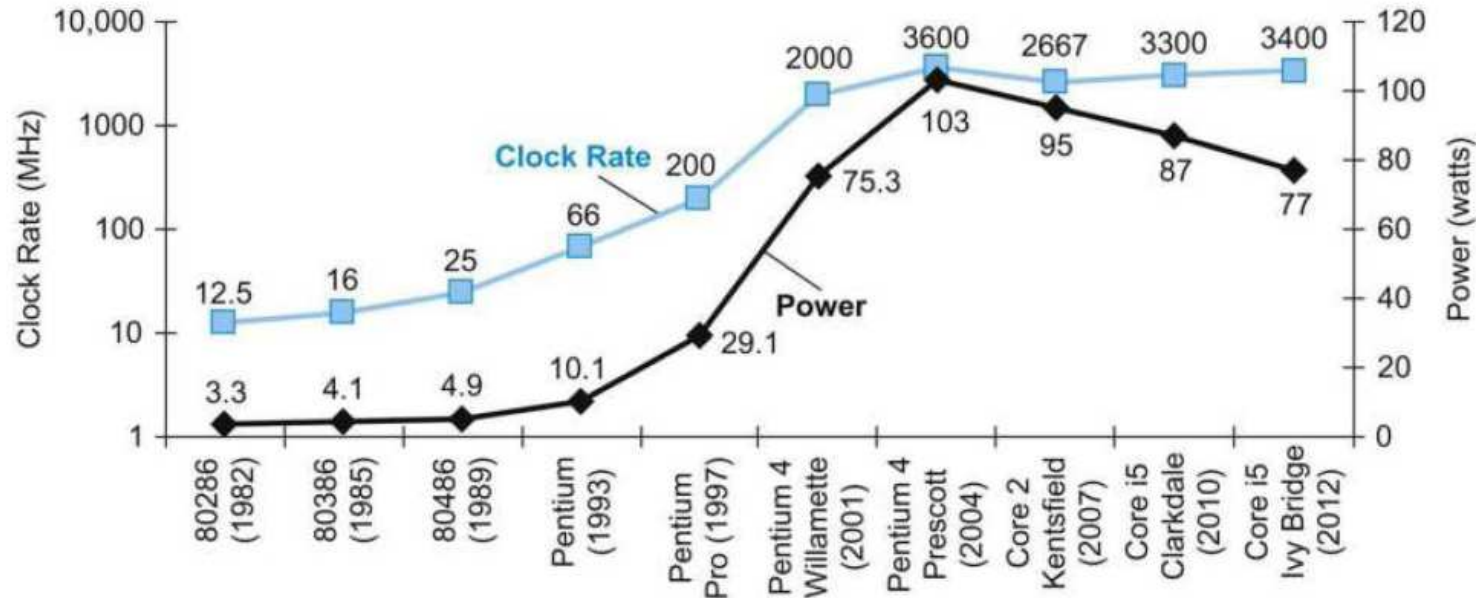
$$\text{ExTime}_{\text{new}} = \text{ExTime}_{\text{old}} \times (0.9 + 0.1/2) = 0.95 \times \text{ExTime}_{\text{old}}$$

$$\text{Speedup}_{\text{overall}} = \frac{\text{ExTime}_{\text{old}}}{0.95 \times \text{ExTime}_{\text{old}}} = 1.053$$

$$\text{Speedup}_{\text{overall}} = \frac{\text{ExTime}_{\text{old}}}{\text{ExTime}_{\text{new}}}$$

# Power wall

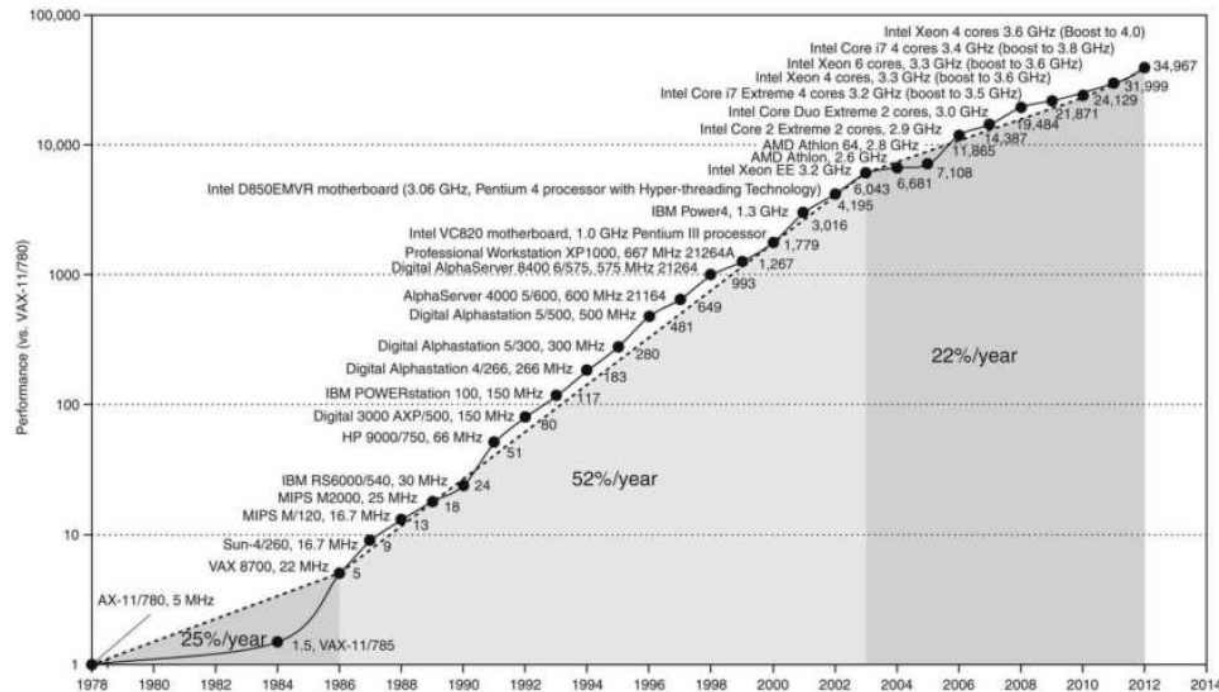
- Power consumption is getting more and more important!



- Cost
- Thermal – Deeply related to performance
- The main reason we are using multi-core CPUs
  - Add more cores, reduce clock frequency, and extract parallelism from software
  - Still hard to exploit full performance available from CPU, why?

# Trends in performance improvements

- Improvement rate per year is saturating



- Limits of power
- Limits in Instruction-Level Parallelism
- Memory wall
- Limits in improving performance of general purpose CPU
  - A rise of domain-specific architecture: e.g., GPU, NPU
  - Nor for a wide range of workloads but for specific workloads